

# Lending Club Default Analysis

The analysis is divided into four main parts:

1. Data understanding
2. Data cleaning (cleaning missing values, removing redundant columns etc.)
3. Data Analysis
4. Recommendations

In [ ]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
loan = pd.read_csv("loan.csv", sep=",")
loan.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 39717 entries, 0 to 39716
```

```
Columns: 111 entries, id to total_il_high_credit_limit
```

```
dtypes: float64(74), int64(13), object(24)
```

```
memory usage: 33.6+ MB
```

```
/Library/Frameworks/Python.framework/Versions/3.5/lib/python3.5/site-packages/IPython/core/interactiveshell.py:2728: DtypeWarning: Columns (47) have mixed types. Specify dtype option on import or set low_memory=False.
```

```
interactivity=interactivity, compiler=compiler, result=result)
```

## Data Understanding


In [ ]:

```
# Let's look at the first few rows of the df
loan.head()
```

Out[ ]:

	id	member_id	loan_amnt	funded_amnt	funded_amnt_inv	term	int_rate	installment	grade
0	1077501	1296599	5000	5000	4975.0	36 months	10.65%	162.87	F
1	1077430	1314167	2500	2500	2500.0	60 months	15.27%	59.83	C
2	1077175	1313524	2400	2400	2400.0	36 months	15.96%	84.33	C
3	1076863	1277178	10000	10000	10000.0	36 months	13.49%	339.31	C
4	1075358	1311748	3000	3000	3000.0	60 months	12.69%	67.79	F

5 rows × 111 columns



```
In [ ]: # Looking at all the column names
        loan.columns
```

```
Out[ ]: Index(['id', 'member_id', 'loan_amnt', 'funded_amnt', 'funded_amnt_inv',
              'term', 'int_rate', 'installment', 'grade', 'sub_grade',
              ...,
              'num_tl_90g_dpd_24m', 'num_tl_op_past_12m', 'pct_tl_nvr_dlq',
              'percent_bc_gt_75', 'pub_rec_bankruptcies', 'tax_liens',
              'tot_hi_cred_lim', 'total_bal_ex_mort', 'total_bc_limit',
              'total_il_high_credit_limit'],
              dtype='object', length=111)
```

Some of the important columns in the dataset are loan\_amount, term, interest rate, grade, sub grade, annual income, purpose of the loan etc.

The **target variable**, which we want to compare across the independent variables, is loan status. The strategy is to figure out compare the average default rates across various independent variables and identify the ones that affect default rate the most.

## Data Cleaning

Some columns have a large number of missing values, let's first fix the missing values and then check for other types of data quality problems.

```
In [ ]: # summarising number of missing values in each column
        loan.isnull().sum()
```

```
Out[ ]: id                                0
        member_id                        0
        loan_amnt                        0
        funded_amnt                      0
        funded_amnt_inv                  0
        term                             0
        int_rate                         0
        installment                      0
        grade                            0
        sub_grade                        0
        emp_title                        2459
        emp_length                       1075
        home_ownership                    0
        annual_inc                        0
        verification_status               0
        issue_d                           0
        loan_status                       0
        pymnt_plan                        0
        url                               0
        desc                             12940
        purpose                           0
        title                             11
        zip_code                          0
        addr_state                        0
        dti                               0
        delinq_2yrs                       0
        earliest_cr_line                  0
```

```

inq_last_6mths      0
mths_since_last_delinq  25682
mths_since_last_record  36931
...
mo_sin_old_rev_tl_op  39717
mo_sin_rcnt_rev_tl_op  39717
mo_sin_rcnt_tl        39717
mort_acc              39717
mths_since_recent_bc  39717
mths_since_recent_bc_dlq  39717
mths_since_recent_inq  39717
mths_since_recent_revol_delinq  39717
num_accts_ever_120_pd  39717
num_actv_bc_tl        39717
num_actv_rev_tl        39717
num_bc_sats            39717
num_bc_tl              39717
num_il_tl              39717
num_op_rev_tl          39717
num_rev_accts          39717
num_rev_tl_bal_gt_0    39717
num_sats               39717
num_tl_120dpd_2m       39717
num_tl_30dpd           39717
num_tl_90g_dpd_24m     39717
num_tl_op_past_12m     39717
pct_tl_nvr_dlq         39717
percent_bc_gt_75       39717
pub_rec_bankruptcies   697
tax_liens              39
tot_hi_cred_lim         39717
total_bal_ex_mort       39717
total_bc_limit          39717
total_il_high_credit_limit  39717
Length: 111, dtype: int64

```

```

In [ ]: # percentage of missing values in each column
        round(loan.isnull().sum()/len(loan.index), 2)*100

```

```

Out[ ]: id              0.0
        member_id       0.0
        loan_amnt       0.0
        funded_amnt     0.0
        funded_amnt_inv  0.0
        term            0.0
        int_rate        0.0
        installment     0.0
        grade           0.0
        sub_grade       0.0
        emp_title       6.0
        emp_length      3.0
        home_ownership  0.0
        annual_inc      0.0
        verification_status  0.0
        issue_d         0.0
        loan_status     0.0
        pymnt_plan      0.0
        url             0.0
        desc            33.0

```

```

purpose      0.0
title        0.0
zip_code     0.0
addr_state   0.0
dti          0.0
delinq_2yrs  0.0
earliest_cr_line 0.0
inq_last_6mths 0.0
mths_since_last_delinq 65.0
mths_since_last_record 93.0
...
mo_sin_old_rev_tl_op 100.0
mo_sin_rcnt_rev_tl_op 100.0
mo_sin_rcnt_tl 100.0
mort_acc 100.0
mths_since_recent_bc 100.0
mths_since_recent_bc_dlq 100.0
mths_since_recent_inq 100.0
mths_since_recent_revol_delinq 100.0
num_accts_ever_120_pd 100.0
num_actv_bc_tl 100.0
num_actv_rev_tl 100.0
num_bc_sats 100.0
num_bc_tl 100.0
num_il_tl 100.0
num_op_rev_tl 100.0
num_rev_accts 100.0
num_rev_tl_bal_gt_0 100.0
num_sats 100.0
num_tl_120dpd_2m 100.0
num_tl_30dpd 100.0
num_tl_90g_dpd_24m 100.0
num_tl_op_past_12m 100.0
pct_tl_nvr_dlq 100.0
percent_bc_gt_75 100.0
pub_rec_bankruptcies 2.0
tax_liens 0.0
tot_hi_cred_lim 100.0
total_bal_ex_mort 100.0
total_bc_limit 100.0
total_il_high_credit_limit 100.0
Length: 111, dtype: float64

```

You can see that many columns have 100% missing values, some have 65%, 33% etc. First, let's get rid of the columns having 100% missing values.

```

In [ ]: # removing the columns having more than 90% missing values
missing_columns = loan.columns[100*(loan.isnull().sum()/len(loan.index)) > 90]
print(missing_columns)

```

```

Index(['mths_since_last_record', 'next_pymnt_d', 'mths_since_last_major_derog',
      'annual_inc_joint', 'dti_joint', 'verification_status_joint',
      'tot_coll_amt', 'tot_cur_bal', 'open_acc_6m', 'open_il_6m',
      'open_il_12m', 'open_il_24m', 'mths_since_rcnt_il', 'total_bal_il',
      'il_util', 'open_rv_12m', 'open_rv_24m', 'max_bal_bc', 'all_util',
      'total_rev_hi_lim', 'inq_fi', 'total_cu_tl', 'inq_last_12m',
      'acc_open_past_24mths', 'avg_cur_bal', 'bc_open_to_buy', 'bc_util',
      'mo_sin_old_il_acct', 'mo_sin_old_rev_tl_op', 'mo_sin_rcnt_rev_tl_op',
      'mo_sin_rcnt_tl', 'mort_acc', 'mths_since_recent_bc',

```

```
'mths_since_recent_bc_dlq', 'mths_since_recent_inq',
'mths_since_recent_revol_delinq', 'num_accts_ever_120_pd',
'num_actv_bc_tl', 'num_actv_rev_tl', 'num_bc_sats', 'num_bc_tl',
'num_il_tl', 'num_op_rev_tl', 'num_rev_accts', 'num_rev_tl_bal_gt_0',
'num_sats', 'num_tl_120dpd_2m', 'num_tl_30dpd', 'num_tl_90g_dpd_24m',
'num_tl_op_past_12m', 'pct_tl_nvr_dlq', 'percent_bc_gt_75',
'tot_hi_cred_lim', 'total_bal_ex_mort', 'total_bc_limit',
'total_il_high_credit_limit'],
dtype='object')
```

```
In [ ]: loan = loan.drop(missing_columns, axis=1)
print(loan.shape)
```

```
(39717, 55)
```

```
In [ ]: # summarise number of missing values again
100*(loan.isnull().sum()/len(loan.index))
```

```
Out[ ]: id                                0.000000
member_id                             0.000000
loan_amnt                             0.000000
funded_amnt                           0.000000
funded_amnt_inv                       0.000000
term                                  0.000000
int_rate                              0.000000
installment                           0.000000
grade                                 0.000000
sub_grade                             0.000000
emp_title                             6.191303
emp_length                             2.706650
home_ownership                        0.000000
annual_inc                             0.000000
verification_status                  0.000000
issue_d                               0.000000
loan_status                           0.000000
pymnt_plan                           0.000000
url                                   0.000000
desc                                 32.580507
purpose                               0.000000
title                                 0.027696
zip_code                             0.000000
addr_state                            0.000000
dti                                   0.000000
delinq_2yrs                           0.000000
earliest_cr_line                       0.000000
inq_last_6mths                         0.000000
mths_since_last_delinq                 64.662487
open_acc                               0.000000
pub_rec                               0.000000
revol_bal                             0.000000
revol_util                             0.125891
total_acc                             0.000000
initial_list_status                    0.000000
out_prncp                             0.000000
out_prncp_inv                         0.000000
total_pymnt                           0.000000
total_pymnt_inv                       0.000000
total_rec_prncp                       0.000000
total_rec_int                         0.000000
```

```
total_rec_late_fee      0.000000
recoveries              0.000000
collection_recovery_fee 0.000000
last_pymnt_d           0.178765
last_pymnt_amnt         0.000000
last_credit_pull_d      0.005036
collections_12_mths_ex_med 0.140998
policy_code             0.000000
application_type        0.000000
acc_now_delinq          0.000000
chargeoff_within_12_mths 0.140998
delinq_amnt             0.000000
pub_rec_bankruptcies    1.754916
tax_liens               0.098195
dtype: float64
```

```
In [ ]: # There are now 2 columns having approx 32 and 64% missing values -
        # description and months since last delinquent

        # Let's have a look at a few entries in the columns
        loan.loc[:, ['desc', 'mths_since_last_delinq']].head()
```

```
Out[ ]:
```

	desc	mths_since_last_delinq
0	Borrower added on 12/22/11 > I need to upgra...	NaN
1	Borrower added on 12/22/11 > I plan to use t...	NaN
2	NaN	NaN
3	Borrower added on 12/21/11 > to pay for prop...	35.0
4	Borrower added on 12/21/11 > I plan on combi...	38.0

The column description contains the comments the applicant had written while applying for the loan. Although one can use some text analysis techniques to derive new features from this column (such as sentiment, number of positive/negative words etc.), we will not use this column in this analysis.

Secondly, months since last delinquent represents the number months passed since the person last fell into the 90 DPD group. There is an important reason we shouldn't use this column in analysis - since at the time of loan application, we will not have this data (it gets generated months after the loan has been approved), it cannot be used as a predictor of default at the time of loan approval.

Thus let's drop the two columns.

```
In [ ]: # dropping the two columns
        loan = loan.drop(['desc', 'mths_since_last_delinq'], axis=1)
```

```
In [ ]: # summarise number of missing values again
        100*(loan.isnull().sum()/len(loan.index))
```

```
Out[ ]:
```

id	0.000000
member_id	0.000000

loan_amnt	0.000000
funded_amnt	0.000000
funded_amnt_inv	0.000000
term	0.000000
int_rate	0.000000
installment	0.000000
grade	0.000000
sub_grade	0.000000
emp_title	6.191303
emp_length	2.706650
home_ownership	0.000000
annual_inc	0.000000
verification_status	0.000000
issue_d	0.000000
loan_status	0.000000
pymnt_plan	0.000000
url	0.000000
purpose	0.000000
title	0.027696
zip_code	0.000000
addr_state	0.000000
dti	0.000000
delinq_2yrs	0.000000
earliest_cr_line	0.000000
inq_last_6mths	0.000000
open_acc	0.000000
pub_rec	0.000000
revol_bal	0.000000
revol_util	0.125891
total_acc	0.000000
initial_list_status	0.000000
out_prncp	0.000000
out_prncp_inv	0.000000
total_pymnt	0.000000
total_pymnt_inv	0.000000
total_rec_prncp	0.000000
total_rec_int	0.000000
total_rec_late_fee	0.000000
recoveries	0.000000
collection_recovery_fee	0.000000
last_pymnt_d	0.178765
last_pymnt_amnt	0.000000
last_credit_pull_d	0.005036
collections_12_mths_ex_med	0.140998
policy_code	0.000000
application_type	0.000000
acc_now_delinq	0.000000
chargeoff_within_12_mths	0.140998
delinq_amnt	0.000000
pub_rec_bankruptcies	1.754916
tax_liens	0.098195
dtype:	float64

There are some more columns with missing values, but let's ignore them for now (since we are not doing any modeling, we don't need to impute all missing values anyway).

But let's check whether some rows have a large number of missing values.

```
In [ ]: # missing values in rows  
        loan.isnull().sum(axis=1)
```

```
Out[ ]: 0      1  
        1      0  
        2      1  
        3      0  
        4      0  
        5      0  
        6      0  
        7      0  
        8      1  
        9      0  
       10      0  
       11      0  
       12      0  
       13      0  
       14      0  
       15      0  
       16      0  
       17      0  
       18      0  
       19      0  
       20      0  
       21      0  
       22      0  
       23      0  
       24      0  
       25      0  
       26      1  
       27      0  
       28      0  
       29      0  
       ..  
    39687      4  
    39688      4  
    39689      4  
    39690      4  
    39691      4  
    39692      4  
    39693      4  
    39694      4  
    39695      4  
    39696      4  
    39697      4  
    39698      4  
    39699      4  
    39700      5  
    39701      4  
    39702      4  
    39703      4  
    39704      5  
    39705      4  
    39706      5  
    39707      4  
    39708      4  
    39709      4  
    39710      4
```



```

39711    4
39712    4
39713    4
39714    5
39715    5
39716    4
Length: 39717, dtype: int64

```

```

In [ ]: # checking whether some rows have more than 5 missing values
        len(loan[loan.isnull().sum(axis=1) > 5].index)

```

```

Out[ ]: 0

```

The data looks clean by and large. Let's also check whether all columns are in the correct format.

```

In [ ]: loan.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 39717 entries, 0 to 39716
Data columns (total 53 columns):
id                39717 non-null int64
member_id         39717 non-null int64
loan_amnt         39717 non-null int64
funded_amnt       39717 non-null int64
funded_amnt_inv   39717 non-null float64
term              39717 non-null object
int_rate          39717 non-null object
installment       39717 non-null float64
grade             39717 non-null object
sub_grade         39717 non-null object
emp_title         37258 non-null object
emp_length        38642 non-null object
home_ownership    39717 non-null object
annual_inc        39717 non-null float64
verification_status 39717 non-null object
issue_d           39717 non-null object
loan_status       39717 non-null object
pymnt_plan        39717 non-null object
url               39717 non-null object
purpose           39717 non-null object
title             39706 non-null object
zip_code          39717 non-null object
addr_state        39717 non-null object
dti               39717 non-null float64
delinq_2yrs       39717 non-null int64
earliest_cr_line  39717 non-null object
inq_last_6mths    39717 non-null int64
open_acc          39717 non-null int64
pub_rec           39717 non-null int64
revol_bal         39717 non-null int64
revol_util        39667 non-null object
total_acc         39717 non-null int64
initial_list_status 39717 non-null object
out_prncp         39717 non-null float64
out_prncp_inv     39717 non-null float64
total_pymnt       39717 non-null float64
total_pymnt_inv   39717 non-null float64
total_rec_prncp   39717 non-null float64

```

```

total_rec_int          39717 non-null float64
total_rec_late_fee     39717 non-null float64
recoveries            39717 non-null float64
collection_recovery_fee 39717 non-null float64
last_pymnt_d          39646 non-null object
last_pymnt_amnt       39717 non-null float64
last_credit_pull_d    39715 non-null object
collections_12_mths_ex_med 39661 non-null float64
policy_code           39717 non-null int64
application_type      39717 non-null object
acc_now_delinq        39717 non-null int64
chargeoff_within_12_mths 39661 non-null float64
delinq_amnt           39717 non-null int64
pub_rec_bankruptcies  39020 non-null float64
tax_liens             39678 non-null float64
dtypes: float64(18), int64(13), object(22)
memory usage: 16.1+ MB

```

```

In [ ]: # The column int_rate is character type, let's convert it to float
        loan['int_rate'] = loan['int_rate'].apply(lambda x: pd.to_numeric(x.split("%")[0]))

```

```

In [ ]: # checking the data types
        loan.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 39717 entries, 0 to 39716
Data columns (total 53 columns):
id          39717 non-null int64
member_id   39717 non-null int64
loan_amnt   39717 non-null int64
funded_amnt 39717 non-null int64
funded_amnt_inv 39717 non-null float64
term        39717 non-null object
int_rate     39717 non-null float64
installment 39717 non-null float64
grade        39717 non-null object
sub_grade    39717 non-null object
emp_title    37258 non-null object
emp_length   38642 non-null object
home_ownership 39717 non-null object
annual_inc   39717 non-null float64
verification_status 39717 non-null object
issue_d      39717 non-null object
loan_status  39717 non-null object
pymnt_plan   39717 non-null object
url          39717 non-null object
purpose      39717 non-null object
title        39706 non-null object
zip_code     39717 non-null object
addr_state   39717 non-null object
dti          39717 non-null float64
delinq_2yrs  39717 non-null int64
earliest_cr_line 39717 non-null object
inq_last_6mths 39717 non-null int64
open_acc     39717 non-null int64
pub_rec      39717 non-null int64
revol_bal    39717 non-null int64
revol_util   39667 non-null object

```

```

total_acc                39717 non-null int64
initial_list_status      39717 non-null object
out_prncp                39717 non-null float64
out_prncp_inv            39717 non-null float64
total_pymnt              39717 non-null float64
total_pymnt_inv          39717 non-null float64
total_rec_prncp          39717 non-null float64
total_rec_int            39717 non-null float64
total_rec_late_fee       39717 non-null float64
recoveries               39717 non-null float64
collection_recovery_fee  39717 non-null float64
last_pymnt_d             39646 non-null object
last_pymnt_amnt          39717 non-null float64
last_credit_pull_d       39715 non-null object
collections_12_mths_ex_med 39661 non-null float64
policy_code              39717 non-null int64
application_type         39717 non-null object
acc_now_delinq           39717 non-null int64
chargeoff_within_12_mths 39661 non-null float64
delinq_amnt              39717 non-null int64
pub_rec_bankruptcies     39020 non-null float64
tax_liens                 39678 non-null float64
dtypes: float64(19), int64(13), object(21)
memory usage: 16.1+ MB

```

```

In [ ]: # also, Lets extract the numeric part from the variable employment length

# first, Let's drop the missing values from the column (otherwise the regex code below
loan = loan[~loan['emp_length'].isnull()]

# using regular expression to extract numeric values from the string
import re
loan['emp_length'] = loan['emp_length'].apply(lambda x: re.findall('\d+', str(x))[0])

# convert to numeric
loan['emp_length'] = loan['emp_length'].apply(lambda x: pd.to_numeric(x))

```

```

In [ ]: # Looking at type of the columns again
loan.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 38642 entries, 0 to 39716
Data columns (total 53 columns):
id                38642 non-null int64
member_id         38642 non-null int64
loan_amnt         38642 non-null int64
funded_amnt       38642 non-null int64
funded_amnt_inv   38642 non-null float64
term              38642 non-null object
int_rate          38642 non-null float64
installment       38642 non-null float64
grade             38642 non-null object
sub_grade         38642 non-null object
emp_title         37202 non-null object
emp_length        38642 non-null int64
home_ownership    38642 non-null object
annual_inc        38642 non-null float64
verification_status 38642 non-null object

```

```

issue_d          38642 non-null object
loan_status      38642 non-null object
pymnt_plan       38642 non-null object
url              38642 non-null object
purpose          38642 non-null object
title            38632 non-null object
zip_code         38642 non-null object
addr_state       38642 non-null object
dti              38642 non-null float64
delinq_2yrs      38642 non-null int64
earliest_cr_line 38642 non-null object
inq_last_6mths   38642 non-null int64
open_acc         38642 non-null int64
pub_rec          38642 non-null int64
revol_bal        38642 non-null int64
revol_util       38595 non-null object
total_acc        38642 non-null int64
initial_list_status 38642 non-null object
out_prncp        38642 non-null float64
out_prncp_inv     38642 non-null float64
total_pymnt      38642 non-null float64
total_pymnt_inv   38642 non-null float64
total_rec_prncp   38642 non-null float64
total_rec_int     38642 non-null float64
total_rec_late_fee 38642 non-null float64
recoveries       38642 non-null float64
collection_recovery_fee 38642 non-null float64
last_pymnt_d     38576 non-null object
last_pymnt_amnt   38642 non-null float64
last_credit_pull_d 38640 non-null object
collections_12_mths_ex_med 38586 non-null float64
policy_code       38642 non-null int64
application_type  38642 non-null object
acc_now_delinq    38642 non-null int64
chargeoff_within_12_mths 38586 non-null float64
delinq_amnt       38642 non-null int64
pub_rec_bankruptcies 37945 non-null float64
tax_liens         38603 non-null float64
dtypes: float64(19), int64(14), object(20)
memory usage: 15.9+ MB

```

## Data Analysis

Let's now move to data analysis. To start with, let's understand the objective of the analysis clearly and identify the variables that we want to consider for analysis.

The objective is to identify predictors of default so that at the time of loan application, we can use those variables for approval/rejection of the loan. Now, there are broadly three types of variables -

1. those which are related to the applicant (demographic variables such as age, occupation, employment details etc.),
1. loan characteristics (amount of loan, interest rate, purpose of loan etc.) and
2. Customer behaviour variables (those which are generated after the loan is approved such as delinquent 2 years, revolving balance, next payment date etc.).

Now, the customer behaviour variables are not available at the time of loan application, and thus they cannot be used as predictors for credit approval.

Thus, going forward, we will use only the other two types of variables.

```
In [ ]: behaviour_var = [  
    "delinq_2yrs",  
    "earliest_cr_line",  
    "inq_last_6mths",  
    "open_acc",  
    "pub_rec",  
    "revol_bal",  
    "revol_util",  
    "total_acc",  
    "out_prncp",  
    "out_prncp_inv",  
    "total_pymnt",  
    "total_pymnt_inv",  
    "total_rec_prncp",  
    "total_rec_int",  
    "total_rec_late_fee",  
    "recoveries",  
    "collection_recovery_fee",  
    "last_pymnt_d",  
    "last_pymnt_amnt",  
    "last_credit_pull_d",  
    "application_type"]  
behaviour_var
```

```
Out[ ]: ['delinq_2yrs',  
    'earliest_cr_line',  
    'inq_last_6mths',  
    'open_acc',  
    'pub_rec',  
    'revol_bal',  
    'revol_util',  
    'total_acc',  
    'out_prncp',  
    'out_prncp_inv',  
    'total_pymnt',  
    'total_pymnt_inv',  
    'total_rec_prncp',  
    'total_rec_int',  
    'total_rec_late_fee',  
    'recoveries',  
    'collection_recovery_fee',  
    'last_pymnt_d',  
    'last_pymnt_amnt',  
    'last_credit_pull_d',  
    'application_type']
```

```
In [ ]: # Let's now remove the behaviour variables from analysis  
df = loan.drop(behaviour_var, axis=1)  
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 38642 entries, 0 to 39716
```

Data columns (total 32 columns):

```
id                38642 non-null int64
member_id         38642 non-null int64
loan_amnt         38642 non-null int64
funded_amnt       38642 non-null int64
funded_amnt_inv   38642 non-null float64
term              38642 non-null object
int_rate          38642 non-null float64
installment       38642 non-null float64
grade             38642 non-null object
sub_grade         38642 non-null object
emp_title         37202 non-null object
emp_length        38642 non-null int64
home_ownership    38642 non-null object
annual_inc        38642 non-null float64
verification_status 38642 non-null object
issue_d           38642 non-null object
loan_status        38642 non-null object
pymnt_plan        38642 non-null object
url               38642 non-null object
purpose           38642 non-null object
title             38632 non-null object
zip_code          38642 non-null object
addr_state        38642 non-null object
dti               38642 non-null float64
initial_list_status 38642 non-null object
collections_12_mths_ex_med 38586 non-null float64
policy_code       38642 non-null int64
acc_now_delinq    38642 non-null int64
chargeoff_within_12_mths 38586 non-null float64
delinq_amnt       38642 non-null int64
pub_rec_bankruptcies 37945 non-null float64
tax_liens         38603 non-null float64
dtypes: float64(9), int64(8), object(15)
memory usage: 9.7+ MB
```

Typically, variables such as acc\_now\_delinquent, chargeoff within 12 months etc. (which are related to the applicant's past loans) are available from the credit bureau.

```
In [ ]: # also, we will not be able to use the variables zip code, address, state etc.
# the variable 'title' is derived from the variable 'purpose'
# thus let get rid of all these variables as well

df = df.drop(['title', 'url', 'zip_code', 'addr_state'], axis=1)
```

Next, let's have a look at the target variable - loan\_status. We need to relabel the values to a binary form - 0 or 1, 1 indicating that the person has defaulted and 0 otherwise.

```
In [ ]: df['loan_status'] = df['loan_status'].astype('category')
df['loan_status'].value_counts()
```

```
Out[ ]: Fully Paid      32145
Charged Off    5399
Current        1098
Name: loan_status, dtype: int64
```

You can see that fully paid comprises most of the loans. The ones marked 'current' are neither fully paid nor defaulted, so let's get rid of the current loans. Also, let's tag the other two values as 0 or 1.

```
In [ ]: # filtering only fully paid or charged-off
df = df[df['loan_status'] != 'Current']
df['loan_status'] = df['loan_status'].apply(lambda x: 0 if x=='Fully Paid' else 1)

# converting loan_status to integer type
df['loan_status'] = df['loan_status'].apply(lambda x: pd.to_numeric(x))

# summarising the values
df['loan_status'].value_counts()
```

```
Out [ ]: 0    32145
        1     5399
        Name: loan_status, dtype: int64
```

Next, let's start with univariate analysis and then move to bivariate analysis.

## Univariate Analysis

First, let's look at the overall default rate.


```
In [ ]: # default rate
round(np.mean(df['loan_status']), 2)
```

```
Out [ ]: 0.14000000000000001
```

The overall default rate is about 14%.


Let's first visualise the average default rates across categorical variables.

```
In [ ]: # plotting default rates across grade of the loan
sns.barplot(x='grade', y='loan_status', data=df)
plt.show()
```




```
In [ ]: # Lets define a function to plot loan_status across categorical variables
def plot_cat(cat_var):
    sns.barplot(x=cat_var, y='loan_status', data=df)
    plt.show()
```

```
In [ ]: # compare default rates across grade of loan
plot_cat('grade')
```




Clearly, as the grade of loan goes from A to G, the default rate increases. This is expected because the grade is decided by Lending Club based on the riskiness of the loan.

```
In [ ]: # term: 60 months Loans default more than 36 months Loans
plot_cat('term')
```




```
In [ ]: # sub-grade: as expected - A1 is better than A2 better than A3 and so on
plt.figure(figsize=(16, 6))
plot_cat('sub_grade')
```







```
In [ ]: # home ownership: not a great discriminator
plot_cat('home_ownership')
```



```
In [ ]: # verification_status: surprisingly, verified loans default more than not verifiedb
plot_cat('verification_status')
```



```
In [ ]: # purpose: small business loans default the most, then renewable energy and education
plt.figure(figsize=(16, 6))
plot_cat('purpose')
```



```
In [ ]: # Let's also observe the distribution of Loans across years
# first lets convert the year column into datetime and then extract year and month from
df['issue_d'].head()
```

```
Out [ ]: 0    Dec-11
1    Dec-11
2    Dec-11
3    Dec-11
5    Dec-11
Name: issue_d, dtype: object
```

```
In [ ]: from datetime import datetime
df['issue_d'] = df['issue_d'].apply(lambda x: datetime.strptime(x, '%b-%y'))
```

```
In [ ]: # extracting month and year from issue_date
df['month'] = df['issue_d'].apply(lambda x: x.month)
df['year'] = df['issue_d'].apply(lambda x: x.year)
```

```
In [ ]: # Let's first observe the number of Loans granted across years
df.groupby('year').year.count()
```

```
Out [ ]: year
2007      251
2008     1562
2009     4716
2010    11214
2011    19801
Name: year, dtype: int64
```


You can see that the number of loans has increased steadily across years.

```
In [ ]: # number of Loans across months
df.groupby('month').month.count()
```


```
Out[ ]: month
1      2331
2      2278
3      2632
4      2756
5      2838
6      3094
7      3253
8      3321
9      3394
10     3637
11     3890
12     4120
Name: month, dtype: int64
```

Most loans are granted in December, and in general in the latter half of the year.

```
In [ ]: # Lets compare the default rates across years
# the default rate had suddenly increased in 2011, inspite of reducing from 2008 till 2
plot_cat('year')
```




```
In [ ]: # comparing default rates across months: not much variation across months
plt.figure(figsize=(16, 6))
plot_cat('month')
```



Let's now analyse how the default rate varies across continuous variables.

```
In [ ]: # Loan amount: the median loan amount is around 10,000
sns.distplot(df['loan_amnt'])
plt.show()
```



The easiest way to analyse how default rates vary across continuous variables is to bin the variables into discrete categories.

Let's bin the loan amount variable into small, medium, high, very high.


```
In [ ]: # binning loan amount
def loan_amount(n):
    if n < 5000:
        return 'low'
    elif n >= 5000 and n < 15000:
        return 'medium'
    elif n >= 15000 and n < 25000:
        return 'high'
    else:
        return 'very high'

df['loan_amnt'] = df['loan_amnt'].apply(lambda x: loan_amount(x))
```

```
In [ ]: df['loan_amnt'].value_counts()
```


```
Out [ ]: medium      20157
high          7572
low           7095
very high     2720
Name: loan_amnt, dtype: int64
```

```
In [ ]: # Let's compare the default rates across loan amount type
# higher the loan amount, higher the default rate
plot_cat('loan_amnt')
```



```
In [ ]: # Let's also convert funded amount invested to bins
df['funded_amnt_inv'] = df['funded_amnt_inv'].apply(lambda x: loan_amount(x))
```

```
In [ ]: # funded amount invested
plot_cat('funded_amnt_inv')
```




```
In [ ]: # Lets also convert interest rate to low, medium, high
# binning Loan amount
def int_rate(n):
    if n <= 10:
        return 'low'
    elif n > 10 and n <=15:
        return 'medium'
    else:
        return 'high'

df['int_rate'] = df['int_rate'].apply(lambda x: int_rate(x))
```

```
In [ ]: # comparing default rates across rates of interest
# high interest rates default more, as expected
```

```
plot_cat('int_rate')
```




```
In [ ]:
```

```
# debt to income ratio
def dti(n):
    if n <= 10:
        return 'low'
    elif n > 10 and n <=20:
        return 'medium'
    else:
        return 'high'

df['dti'] = df['dti'].apply(lambda x: dti(x))
```

```
In [ ]:
```

```
# comparing default rates across debt to income ratio
# high dti translates into higher default rates, as expected
plot_cat('dti')
```



```
In [ ]:
```

```
# funded amount
def funded_amount(n):
    if n <= 5000:
        return 'low'
```

```


elif n > 5000 and n <=15000:
    return 'medium'
else:
    return 'high'

df['funded_amnt'] = df['funded_amnt'].apply(lambda x: funded_amount(x))

```

In [ ]:

```
plot_cat('funded_amnt')
```



In [ ]:

```

# installment
def installment(n):
    if n <= 200:
        return 'low'
    elif n > 200 and n <=400:
        return 'medium'
    elif n > 400 and n <=600:
        return 'high'
    else:
        return 'very high'

df['installment'] = df['installment'].apply(lambda x: installment(x))


```

In [ ]:

```

# comparing default rates across installment
# the higher the installment amount, the higher the default rate
plot_cat('installment')


```



```
In [ ]: # annual income
def annual_income(n):
    if n <= 50000:
        return 'low'
    elif n > 50000 and n <=100000:
        return 'medium'
    elif n > 100000 and n <=150000:
        return 'high'
    else:
        return 'very high'

df['annual_inc'] = df['annual_inc'].apply(lambda x: annual_income(x))
```

```
In [ ]: # annual income and default rate
# Lower the annual income, higher the default rate
plot_cat('annual_inc')
```



```
In [ ]: # employment length
# first, Let's drop the missing value observations in emp length
df = df[~df['emp_length'].isnull()]


# binning the variable
```



```
def emp_length(n):
    if n <= 1:
        return 'fresher'
    elif n > 1 and n <=3:
        return 'junior'
    elif n > 3 and n <=7:
        return 'senior'
    else:
        return 'expert'

df['emp_length'] = df['emp_length'].apply(lambda x: emp_length(x))
```

```
In [ ]: # emp_length and default rate
# not much of a predictor of default
plot_cat('emp_length')
```




## Segmented Univariate Analysis

We have now compared the default rates across various variables, and some of the important predictors are purpose of the loan, interest rate, annual income, grade etc.

In the credit industry, one of the most important factors affecting default is the purpose of the loan - home loans perform differently than credit cards, credit cards are very different from debt consolidation loans etc.


This comes from business understanding, though let's again have a look at the default rates across the purpose of the loan.

```
In [ ]: # purpose: small business loans default the most, then renewable energy and education
plt.figure(figsize=(16, 6))
plot_cat('purpose')
```



In the upcoming analyses, we will segment the loan applications across the purpose of the loan, since that is a variable affecting many other variables - the type of applicant, interest rate, income, and finally the default rate.

```
In [ ]: # Lets first look at the number of Loans for each type (purpose) of the Loan
# most Loans are debt consolidation (to repay otehr debts), then credit card, major pur
plt.figure(figsize=(16, 6))
sns.countplot(x='purpose', data=df)
plt.show()
```




Let's analyse the top 4 types of loans based on purpose: consolidation, credit card, home improvement and major purchase.

```
In [ ]: # filtering the df for the 4 types of Loans mentioned above
main_purposes = ["credit_card", "debt_consolidation", "home_improvement", "major_purchase"]
df = df[df['purpose'].isin(main_purposes)]
df['purpose'].value_counts()
```


```
Out [ ]: debt_consolidation    17675
credit_card                4899
home_improvement           2785
major_purchase              2080
Name: purpose, dtype: int64
```

```
In [ ]: # plotting number of loans by purpose
sns.countplot(x=df['purpose'])
plt.show()
```



```
In [ ]: # Let's now compare the default rates across two types of categorical variables
# purpose of loan (constant) and another categorical variable (which changes)

plt.figure(figsize=[10, 6])
sns.barplot(x='term', y="loan_status", hue='purpose', data=df)
plt.show()
```




```
In [ ]: # Lets write a function which takes a categorical variable and plots the default rate
# segmented by purpose

def plot_segmented(cat_var):
```


```
plt.figure(figsize=(10, 6))
sns.barplot(x=cat_var, y='loan_status', hue='purpose', data=df)
plt.show()
```

```
plot_segmented('term')
```




In [ ]:

```
# grade of Loan
plot_segmented('grade')
```




```
In [ ]: # home ownership  
plot_segmented('home_ownership')
```




In general, debt consolidation loans have the highest default rates. Lets compare across other categories as well.


```
In [ ]: # year  
plot_segmented('year')
```




```
In [ ]: # emp_length
plot_segmented('emp_length')
```




```
In [ ]: # Loan_amnt: same trend across Loan purposes
plot_segmented('loan_amnt')
```




```
In [ ]: # interest rate
plot_segmented('int_rate')
```




```
In [ ]: # installment
plot_segmented('installment')
```



```
In [ ]: # debt to income ratio
plot_segmented('dti')
```



```
In [ ]: # annual income
plot_segmented('annual_inc')
```



A good way to quantify the effect of a categorical variable on default rate is to see 'how much does the default rate vary across the categories'.

Let's see an example using annual\_inc as the categorical variable.

```
In [ ]: # variation of default rate across annual_inc
df.groupby('annual_inc').loan_status.mean().sort_values(ascending=False)
```



```
Out[ ]: annual_inc
low      0.157966
medium   0.130075
very high 0.101570
high     0.097749
Name: loan_status, dtype: float64
```

```
In [ ]: # one can write a function which takes in a categorical variable and computed the average
# default rate across the categories
# It can also compute the 'difference between the highest and the lowest default rate'
# categories, which is a decent metric indicating the effect of the variable on default

def diff_rate(cat_var):
    default_rates = df.groupby(cat_var).loan_status.mean().sort_values(ascending=False)
    return (round(default_rates, 2), round(default_rates[0] - default_rates[-1], 2))

default_rates, diff = diff_rate('annual_inc')
print(default_rates)
print(diff)

annual_inc
low      0.16
medium   0.13
very high 0.10
high     0.10
Name: loan_status, dtype: float64
0.06
```

Thus, there is a 6% increase in default rate as you go from high to low annual income. We can compute this difference for all the variables and roughly identify the ones that affect default rate the most.

```
In [ ]: # filtering all the object type variables
df_categorical = df.loc[:, df.dtypes == object]
df_categorical['loan_status'] = df['loan_status']

# Now, for each variable, we can compute the incremental diff in default rates
print([i for i in df.columns])

['id', 'member_id', 'loan_amnt', 'funded_amnt', 'funded_amnt_inv', 'term', 'int_rate',
'installment', 'grade', 'sub_grade', 'emp_title', 'emp_length', 'home_ownership', 'annual_inc',
'verification_status', 'issue_d', 'loan_status', 'pymnt_plan', 'purpose', 'dti',
'initial_list_status', 'collections_12_mths_ex_med', 'policy_code', 'acc_now_delinq', 'chargeoff_within_12_mths',
'delinq_amnt', 'pub_rec_bankruptcies', 'tax_liens', 'month', 'year']

/Library/Frameworks/Python.framework/Versions/3.5/lib/python3.5/site-packages/ipykernel_launcher.py:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy
This is separate from the ipykernel package so we can avoid doing imports until
```

```
In [ ]: # storing the diff of default rates for each column in a dict
d = {key: diff_rate(key)[1]*100 for key in df_categorical.columns if key != 'loan_status'}
```

```
print(d)
```

```
{'loan_amnt': 7.0000000000000009, 'funded_amnt_inv': 6.0, 'pymnt_plan': 0.0, 'verification_status': 4.0, 'emp_title': 100.0, 'dti': 5.0, 'home_ownership': 16.0, 'purpose': 5.0, 'sub_grade': 46.0, 'grade': 27.0, 'funded_amnt': 5.0, 'installment': 3.0, 'initial_list_status': 0.0, 'int_rate': 19.0, 'term': 15.0, 'annual_inc': 6.0, 'emp_length': 2.0}
```